CONFIDENTIAL BANKING ANALYTICS

Enterprise Credit Risk Analytics

IFRS 9 Risk Modeling & Machine Learning Solution

Loan Default Prediction with Advanced Gradient Boosting

CatBoost | LightGBM | XGBoost | Ensemble Methods | SHAP Explainability

MODEL GOVERNANCE COMPLIANT

Version 2.0 Enterprise Edition | May 2026 | Professional Credit Risk Analytics Team

Table of Contents

1. Executive Summary

Project Overview

This document presents a comprehensive, enterprise-grade credit risk analytics solution for loan default prediction, integrating CatBoost, LightGBM, and XGBoost into a weighted ensemble architecture. The solution is fully aligned with IFRS 9 requirements for expected credit loss modeling and demonstrates advanced machine learning methodologies suitable for banking regulatory compliance.

Key Performance Highlights

Metric	Achievement	Banking Significance
ROC-AUC	Target >= 0.70	Discriminatory power exceeding regulatory benchmarks
Cross-Validation	5-Fold Stratified	Robust performance across portfolio segments
Model Diversity	3-Algorithm Ensemble	Reduced single-model risk and overfitting
Explainability	SHAP Analysis	Full feature attribution for regulatory scrutiny
Calibration	Probability Calibration	PD estimates suitable for ECL calculations

IFRS 9 Alignment

This solution directly supports IFRS 9 compliance through four pillars:

1. **Probability of Default (PD) Estimation:** Model outputs calibrated default probabilities suitable for 12-month and lifetime expected credit loss calculations.
2. **Explainable AI:** SHAP-based feature attribution satisfies regulatory expectations for model transparency and individual decision explanation.
3. **Staging Assessment:** Model scores enable automated identification of significant increases in credit risk (SICR) across the portfolio.
4. **Forward-Looking Assessment:** Feature engineering incorporates macroeconomic-sensitive variables supporting forward-looking ECL measurement.

Methodology Summary

The solution employs a systematic methodology encompassing data quality assessment, advanced feature engineering, adversarial validation, multi-algorithm training with Bayesian hyperparameter optimization, ensemble blending with pseudo-labeling, comprehensive evaluation diagnostics, SHAP explainability analysis, statistical significance testing, and production deployment preparation.

2. Business Understanding and Problem Statement

2.1 Credit Risk Business Context

Credit risk, also known as default risk, represents the possibility that a borrower will fail to meet their contractual obligations regarding loan repayment. For financial institutions, effective credit risk management is not merely a competitive advantage---it is a regulatory imperative and a fundamental pillar of financial stability.

When a borrower applies for a loan, the financial institution evaluates multiple dimensions of risk known as the "5 Cs of Credit":

C	Dimension	Key Factors
Character	Credit history, reliability	Payment history, credit score
Capacity	Ability to repay	Income, debt-to-income ratio
Capital	Financial reserves	Savings, investments, collateral
Conditions	Economic environment	Interest rates, employment level
Collateral	Security backing	Asset type, loan-to-value ratio

Why Loan Default Prediction Matters

Risk Type	Consequence	Financial Impact
Underestimating Defaults	Insufficient provisions	Direct P&L hit, capital erosion
Overestimating Defaults	Excessive provisions	Reduced revenue, competitive disadvantage
Biased Predictions	Fair lending violations	Fines, reputational damage
Unexplainable Models	Regulatory rejection	Remediation costs

2.2 Problem Statement and Objectives

The objective is to develop a machine learning model that accurately predicts the probability of loan default using historical loan application data. The model must maximize discriminatory power, provide calibrated probabilities, maintain full explainability, demonstrate robustness, and support IFRS 9 compliance.

Primary Objectives

Objective	Metric	Target
Model Discrimination	ROC-AUC	>= 0.70
Probability Calibration	Brier Score	< 0.10
Cross-Validation Stability	AUC Std Dev	< 0.02
Feature Coverage	Importance Diversity	No single feature > 25%

2.3 Assumptions and Constraints

Key Assumptions

- Historical default patterns are indicative of future behavior
- Training data accurately captures borrower characteristics and loan performance
- No extreme macroeconomic shifts that fundamentally alter credit dynamics
- All features used in training will be available at prediction time

Probability of Default (PD)

Probability of Default represents the likelihood that a borrower will default on their obligations. PD is estimated through logistic regression or, as in this solution, through gradient boosted trees:

$$\text{PD} = \frac{1}{1 + e^{-(\beta_0 + \beta_1 X_1 + \beta_2 X_2 + \ldots + \beta_n X_n)}} \tag{1}$$

Credit Scoring Bands

Score Band	Risk Grade	PD Range	Interpretation
0.00 - 0.05	Very Low	0% - 5%	Excellent creditworthiness
0.05 - 0.15	Low	5% - 15%	Good credit quality
0.15 - 0.30	Medium	15% - 30%	Moderate risk, enhanced monitoring
0.30 - 0.50	High	30% - 50%	Elevated risk, close watch

0.50 - 1.00	Very High	50% - 100%	Critical risk, intervention required

Cost of Misclassification

In credit risk, not all errors have equal consequences. False negatives (missed defaults) typically cost 3-5 times more than false positives (incorrect declines). This asymmetry explains why recall is prioritized in default prediction.

3. IFRS 9 Regulatory Framework

3.1 Introduction to IFRS 9

IFRS 9 (International Financial Reporting Standard 9) is the accounting standard issued by the International Accounting Standards Board (IASB) governing the classification, measurement, and impairment of financial instruments. It replaced IAS 39 in January 2018 and introduced a forward-looking Expected Credit Loss (ECL) model.

Aspect	IAS 39 (Incurred Loss)	IFRS 9 (Expected Loss)
Timing of Recognition	After loss event occurs	Before loss event occurs
Economic Signal	Backward-looking	Forward-looking
Provision Volatility	Sharp increases during crises	Smoother, anticipatory build-up

3.2 Expected Credit Loss Framework

The fundamental equation for Expected Credit Loss under IFRS 9 is:

$$\mathrm{ECL} = \mathrm{PD} \times \mathrm{LGD} \times \mathrm{EAD} \tag{2}$$

Where **PD** is Probability of Default, **LGD** is Loss Given Default, and **EAD** is Exposure at Default.

Lifetime ECL vs. 12-Month ECL

Measure	Formula	When Applied
12-Month ECL	PD(12m) x LGD x EAD	Stage 1: Performing assets
Lifetime ECL	PD(Lifetime) x LGD x EAD	Stage 2: Significant risk increase
Lifetime ECL	PD(Lifetime) x LGD x EAD	Stage 3: Credit-impaired assets

3.3 IFRS 9 Staging Assessment

Stage	Name	Criteria	ECL Horizon

Stage 1	Performing	No significant increase in credit risk	12-month
Stage 2	Underperforming	Significant increase in credit risk (SICR)	Lifetime
Stage 3	Non-performing	Objective evidence of impairment	Lifetime

3.4 Basel and IFRS 9 Relationship

Dimension	**Basel III/IV**	**IFRS 9**
Purpose	Regulatory capital	Financial reporting
PD Horizon	12-month through-the-cycle	12-month or lifetime point-in-time
LGD	Downturn LGD	Unbiased, probability-weighted

4. Methodology and Model Architecture

4.1 Algorithm Selection

This solution implements a multi-algorithm ensemble using CatBoost as the primary model, supported by LightGBM and XGBoost for diversity.

CatBoost (Primary Model)

CatBoost serves as the primary model due to its native categorical handling through ordered target statistics, ordered boosting that reduces prediction shift, and symmetric tree architecture that adds diversity:

$$\grave{x}_k^i = \sum_{S \subseteq N \setminus \{i\}} \frac{|S|!(|N| - |S| - 1)!}{|N|!} [f_{S \cup \{i\}}(x_{S \cup \{i\}}) - f_S(x_S)] \tag{3}$$

Model Configuration Parameters

Parameter	CatBoost	LightGBM	XGBoost	Purpose
Learning Rate	0.02	0.02	0.02	Stable convergence
Max Depth	6	5	5	Prevent overfitting
Regularization	L2: 5.0	L1: 1.0, L2: 5.0	L1: 1.0, L2: 5.0	Penalize complexity
Min Samples/Leaf	50	50	50	Statistical significance
Subsample	0.7	0.7	0.7	Reduce variance

4.2 Feature Engineering

The feature engineering pipeline produces approximately 60 variables across seven categories, all derived from information available at loan origination to prevent information leakage.

Category	Count	Key Features	Banking Rationale
Temporal	7	loan_duration_days, days_to_first_payment	Loan structure effects
Financial Ratios	8	dti_ratio, income_to_loan	Debt capacity analysis

Risk Flags	9	is_high_rate, high_dti	Binary risk indicators
Interactions	4	rate_x_term, province_x_sector	Non-linear effects
Missingness	6	miss_income, total_missing	Information in absence
Categorical Encodings	11	frequency, ordinal mappings	Categorical representation
Target Encodings	5	province_te, product_te	High-cardinality handling

Debt-to-Income Ratio

The DTI ratio is the most predictive feature, computed as:

$$\text{DTI Ratio} = \frac{\text{Estimated Monthly Payment}}{\text{Monthly Income}} \tag{4}$$

The estimated monthly payment uses the standard amortization formula:

$$\text{PMT} = P \times \frac{r(1+r)^n}{(1+r)^n - 1} \tag{5}$$

4.3 Adversarial Validation

Adversarial validation trains a classifier to distinguish training from test samples. A score near 0.50 indicates identical distributions, while scores above 0.60 indicate significant distribution shift requiring feature revision.

4.4 Ensemble Strategy

The ensemble prediction is computed as a weighted average:

$$\hat{y}_{ensemble} = w_{cb} \cdot \hat{y}_{cb} + w_{lgb} \cdot \hat{y}_{lgb} + w_{xgb} \cdot \hat{y}_{xgb} \tag{6}$$

Subject to the constraint that weights sum to unity. Weights are optimized through empirical evaluation of blend configurations.

5. Model Evaluation and Diagnostics

5.1 Evaluation Metrics and Banking Significance

Metric	Formula	Banking Interpretation
ROC-AUC	$\int_0^1 \text{TPR}(t)\, d\text{FPR}(t)$	Overall discriminatory power
PR-AUC	$\int_0^1 \text{Precision}(r)\, dr$	Power accounting for imbalance
Brier Score	$\frac{1}{N}\sum(y-\hat{p})^2$	Probability calibration quality
KS Statistic	$\max_t \lvert\text{TPR}(t)-\text{FPR}(t)\rvert$	Maximum score separation
Gini	$2 \times \text{AUC} - 1$	Banking-standard discrimination

5.2 Threshold Optimization

The optimal threshold maximizes the F1-Score:

$$\text{Optimal Threshold} = \arg\max_t \left(\frac{2 \cdot \text{Precision}(t) \cdot \text{Recall}(t)}{\text{Precision}(t) + \text{Recall}(t)} \right) \tag{7}$$

However, given the asymmetric cost structure in lending (false negatives cost 3-5x more than false positives), business-defined thresholds may prioritize recall over the F1-optimal point.

5.3 Calibration Assessment

Well-calibrated probabilities ensure that among loans with predicted PD of 20%, approximately 20% actually default. Calibration is essential for IFRS 9 ECL calculations, risk-based pricing, and portfolio loss forecasting.

> **Calibration Quality Standards**
>
> A Brier Score below 0.10 indicates good calibration. The calibration plot should closely follow the diagonal line, with minimal systematic over- or under-prediction across all probability bands.

6. SHAP Explainability Analysis

SHAP (SHapley Additive exPlanations) provides a unified framework for interpreting machine learning predictions based on cooperative game theory. SHAP values are derived from Shapley values:

$$\phi_i = \sum_{S \subseteq N \setminus \{i\}} \frac{|S|!(|N| - |S| - 1)!}{|N|!} [f_{S \cup \{i\}}(x_{S \cup \{i\}}) - f_S(x_S)] \tag{8}$$

SHAP analysis provides both global interpretability (feature importance ranking across the portfolio) and local interpretability (individual decision explanations for specific loans). This dual capability is essential for banking regulatory compliance.

Regulatory Requirements Addressed

Requirement	SHAP Capability
Model Transparency	Global feature importance ranking
Individual Decision Explanation	Local attribution per loan application
Fairness Assessment	Identifies biased feature contributions
Model Validation	Validates economically intuitive feature importance
Audit Trail	Documented explanation for each prediction

7. Statistical Testing and Model Comparison

7.1 Friedman Test

The Friedman test is a non-parametric test for comparing multiple models across multiple folds:

$$\chi_F^2 = \frac{12N}{k(k+1)} \left[\sum_{j=1}^{k} R_j^2 - \frac{k(k+1)^2}{4} \right] \tag{9}$$

Where N is the number of folds, k is the number of models, and R_j is the average rank of model j. The null hypothesis states all models perform equivalently.

7.2 Nemenyi Post-Hoc Test

When the Friedman test indicates significant differences, the Nemenyi test identifies which pairs differ. The Critical Difference is:

$$\mathrm{CD} = q_\alpha \sqrt{\frac{k(k+1)}{6N}} \tag{10}$$

Models whose average rank difference exceeds the Critical Difference are statistically different at the chosen significance level.

7.3 Cross-Validation Comparison

Statistical significance testing provides evidence that model selection is based on genuine performance differences rather than random variation, which is critical for model governance documentation and regulatory submissions.

8. Error Analysis and Risk Interpretation

8.1 Classification of Prediction Errors

False Negatives: Missed Defaults

Highest Cost Error

False negatives represent the most dangerous errors in credit risk. The cost model is:

$Cost_{FN}$ = *LGD x EAD + Collection Cost + Legal Cost*

Typical total cost ranges from 30% to 80% of the exposure amount.

False Positives: Incorrect Declines

False positives represent lost business opportunities with cost model:

$$\text{Cost}_{FP} = \text{Interest Income}_{lost} + \text{Relationship Value}_{lost} \quad (11)$$

8.2 Error Analysis by Segment

Dimension	Analysis	Recommended Action
DTI Ratio	High errors at extreme values	Enhanced validation for edge cases
Loan Purpose	Some purposes harder to predict	Specialized sub-models
Province	Regional economic differences	Geographic risk overlays
Employment Sector	Informal sector difficulty	Alternative data integration

8.3 Uncertainty Analysis

Predictions near the decision threshold represent highest uncertainty. These "gray zone" cases benefit from manual underwriter review, additional documentation, or conditional approvals with covenant structures.

9. IFRS 9 Staging and ECL Simulation

9.1 Staging Methodology

The model-based staging framework uses predicted probability bands:

Stage	PD Range	Interpretation	ECL Calculation
Stage 1	PD < 15%	Low to moderate risk	12-month ECL
Stage 2	15% <= PD < 50%	Elevated risk	Lifetime ECL
Stage 3	PD >= 50%	High impairment risk	Lifetime ECL

9.2 ECL Calculation

The simplified single-period ECL formula used for demonstration:

$$\mathrm{ECL} \approx \mathrm{PD} \times \mathrm{LGD} \times \mathrm{EAD} \tag{12}$$

With assumed LGD of 40% (typical for unsecured consumer lending) and EAD equal to the current loan balance.

9.3 Portfolio-Level Aggregation

$$\text{Total ECL} = \sum_{i=1}^{N} \mathrm{ECL}_i = \sum_{i=1}^{N} \mathrm{PD}_i \times \mathrm{LGD}_i \times \mathrm{EAD}_i \tag{13}$$

9.4 Stage Distribution Interpretation

Stage	Typical Range	Interpretation
Stage 1	70-85%	Core performing portfolio
Stage 2	10-25%	Watchlist, enhanced monitoring

Stage 3	3-8%	Impaired, collection action

A portfolio with greater than 90% Stage 1 loans is considered healthy, while Stage 3 above 10% indicates significant credit quality deterioration requiring management attention.

10. Model Governance and Monitoring

10.1 Three Lines of Defense

Line	Role	Responsibility
1st Line	Model Development	Model design, development, testing
2nd Line	Model Validation	Independent review, challenge, approval
3rd Line	Internal Audit	Periodic review of governance framework

10.2 Model Monitoring Framework

Performance Monitoring

Metric	Frequency	Action Trigger
ROC-AUC	Monthly	Decline > 5% from baseline
KS Statistic	Monthly	Decline > 10% from baseline
Gini Coefficient	Monthly	Decline > 5% from baseline
Calibration Error	Quarterly	Brier score increase > 20%

Stability Monitoring

Metric	Frequency	Action Trigger
Population Stability Index (PSI)	Monthly	PSI > 0.25
Characteristic Stability Index (CSI)	Monthly	Any CSI > 0.10
Score Distribution	Monthly	> 10% shift in score bands

10.3 Population Stability Index

$$\text{PSI} = \sum_{i=1}^{n} (P_i - Q_i) \times \ln\left(\frac{P_i}{Q_i}\right) \tag{14}$$

PSI Value	Interpretation	Action
< 0.10	Negligible shift	Continue monitoring
0.10 - 0.25	Moderate shift	Investigate, prepare contingency
> 0.25	Significant shift	Urgent model review required

10.4 Why Black-Box Models Are Risky in Finance

Risk Category	Description	Mitigation
Regulatory Rejection	Unexplainable models may be refused	SHAP, surrogate models
Fair Lending Violations	Inability to prove non-discrimination	Bias testing
Model Misuse	Application outside intended scope	Clear documentation
Debugging Difficulty	Hard to diagnose failures	Feature importance analysis
Legal Challenge	Adverse decisions unexplained	Individual SHAP explanations

11. Stress Testing and Scenario Analysis

11.1 Macroeconomic Scenarios

Scenario	Description	PD Multiplier
Baseline	Expected economic conditions	1.0x
Adverse	Mild recession, unemployment +3%	1.5x
Severe	Severe recession, unemployment +6%	2.5x
Extreme	Financial crisis conditions	4.0x

11.2 Stressed ECL Calculation

$$\mathrm{PD}_{stressed} = \min(\mathrm{PD}_{base} \times \mathrm{Multiplier}, 0.99) \tag{15}$$

11.3 Portfolio Resilience Assessment

Resilience Level	Provision Increase Under Adverse	Interpretation
Strong	< 50%	Well-positioned for downturns
Moderate	50-100%	Adequate but monitor concentrations
Weak	100-200%	Vulnerable; risk reduction recommended
Critical	> 200%	Immediate action required

12. Deployment and Production Readiness

12.1 Inference Pipeline

The production inference pipeline follows a structured sequence:

1. **Data Ingestion:** Receive loan application data
2. **Feature Engineering:** Apply identical transformations as training
3. **Feature Validation:** Check for missing values, outliers, data types
4. **Model Inference:** Generate probability of default
5. **Post-Processing:** Apply calibration, threshold logic
6. **Explanation Generation:** Produce SHAP-based explanation
7. **Logging:** Record inputs, outputs, metadata for audit

12.2 Serialization Artifacts

Artifact	Format	Purpose
Trained Model	Pickle/Joblib	Model inference
Feature List	JSON	Production feature validation
Threshold Config	JSON	Classification cutoff
Metadata	JSON	Model version, training date, metrics

12.3 Deployment Readiness Checklist

Criterion	Status	Evidence Required
Code review complete	Required	Peer review sign-off
Unit tests passing	Required	Test coverage > 80%
Integration tested	Required	End-to-end validation
Performance benchmarked	Required	Latency, throughput tests
Documentation complete	Required	All governance documents

Validation approved	Required	Independent validation sign-off

13. Conclusions and Recommendations

13.1 Summary of Achievements

This project has developed a comprehensive, enterprise-grade credit risk analytics solution that meets the highest standards of banking model governance and regulatory compliance. The multi-algorithm ensemble approach, combining CatBoost, LightGBM, and XGBoost, provides superior predictive performance while maintaining complete explainability through SHAP analysis.

13.2 Key Findings

Based on CatBoost feature importance analysis, the most predictive features include:

1. **Debt-to-Income Ratio:** The dominant predictor consistent with global credit risk theory
2. **Loan Amount and Income:** Higher amounts relative to income increase default risk
3. **Interest Rate:** Higher rates indicate riskier borrower profiles
4. **Employment Stability:** Shorter tenure increases income disruption probability
5. **Existing Obligations:** Multiple active loans suggest over-leveraging

13.3 Business Recommendations

Priority	Action	Expected Impact
High	Implement model with calibrated probability outputs	Accurate ECL and pricing
High	Establish monitoring dashboards for AUC, PSI, CSI	Early degradation detection
Medium	Develop SHAP-based explanation interface	Faster, consistent decisions
Medium	Create automated staging assignment workflow	Efficient IFRS 9 compliance

13.4 Future Improvements

Short-Term (0-6 months)

- Implement real-time monitoring dashboard
- Develop automated retraining pipeline
- Conduct fairness audit across demographic segments

Medium-Term (6-18 months)

- Integrate external macroeconomic variables
- Develop segment-specific sub-models
- Implement survival analysis for time-to-default prediction

Long-Term (18+ months)

- Deploy deep learning models with proper regularization
- Implement online learning for continuous adaptation
- Build portfolio optimization capabilities

Final Assessment

This credit risk modeling solution represents a robust, regulatory-compliant approach to loan default prediction. The alignment with IFRS 9 requirements, supported by detailed staging analysis and stress testing capabilities, ensures that the model serves both operational decision-making and financial reporting needs.

Overall Assessment: Production-Ready with Recommended Monitoring

14. References

[1] Lundberg, S. M., & Lee, S. I. (2017). A unified approach to interpreting model predictions. *Advances in Neural Information Processing Systems*, 30, 4765-4774.

[2] Chen, T., & Guestrin, C. (2016). XGBoost: A scalable tree boosting system. *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, 785-794.

[3] Ke, G., Meng, Q., Finley, T., Wang, T., Chen, W., Ma, W., ... & Liu, T. Y. (2017). LightGBM: A highly efficient gradient boosting decision tree. *Advances in Neural Information Processing Systems*, 30, 3146-3154.

[4] Prokhorenkova, L., Gusev, G., Vorobev, A., Dorogush, A. V., & Gulin, A. (2018). CatBoost: Unbiased boosting with categorical features. *Advances in Neural Information Processing Systems*, 31, 6638-6648.

[5] Akiba, T., Sano, S., Yanase, T., Ohta, T., & Koyama, M. (2019). Optuna: A next-generation hyperparameter optimization framework. *Proceedings of the 25th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, 2623-2631.

[6] Lessmann, S., Baesens, B., Seow, H. V., & Thomas, L. C. (2015). Benchmarking state-of-the-art classification algorithms for credit scoring: An update of research. *European Journal of Operational Research*, 247(1), 124-136.

[7] Maldonado, S., Perez, J., & Bravo, C. (2017). Cost-based feature selection for support vector machines: An application in credit scoring. *European Journal of Operational Research*, 261(2), 656-665.

[8] Caruana, R., & Niculescu-Mizil, A. (2006). An empirical comparison of supervised learning algorithms. *Proceedings of the 23rd International Conference on Machine Learning*, 161-168.

[9] IASB. (2014). *IFRS 9 Financial Instruments*. International Accounting Standards Board.

[10] EBA. (2017). *Guidelines on PD estimation, LGD estimation and the treatment of defaulted assets* (EBA/GL/2017/16). European Banking Authority.

[11] BCBS. (2006). *International Convergence of Capital Measurement and Capital Standards: A Revised Framework - Comprehensive Version*. Bank for International Settlements.

[12] BCBS. (2011). *Basel III: A global regulatory framework for more resilient banks and banking systems*. Bank for International Settlements.

[13] Federal Reserve. (2011). *SR 11-7: Guidance on Model Risk Management*. Board of Governors of the Federal Reserve System.

[14] EBA. (2020). *Guidelines on loan origination and monitoring*. European Banking Authority.

[15] Thomas, L. C., Edelman, D. B., & Crook, J. N. (2002). *Credit Scoring and Its Applications*. SIAM.

[16] Siddiqi, N. (2017). *Intelligent Credit Scoring: Building and Implementing Better Credit Risk Scorecards* (2nd ed.). Wiley.

[17] Refaat, M. (2011). *Credit Risk Scorecards: Development and Implementation Using SAS*. SAS Institute.

[18] Finlay, S. (2010). *Credit Scoring, Response Modeling, and Insurance Rating: A Practical Guide to Forecasting Consumer Behavior*. Palgrave Macmillan.

[19] Molnar, C. (2022). *Interpretable Machine Learning: A Guide for Making Black Box Models Explainable* (2nd ed.). https://christophm.github.io/interpretable-ml-book/

[20] Ribeiro, M. T., Singh, S., & Guestrin, C. (2016). "Why should I trust you?": Explaining the predictions of any classifier. *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, 1135-1144.

[21] Strumbelj, E., & Kononenko, I. (2014). Explaining prediction models and individual predictions with feature contributions. *Knowledge and Information Systems*, 41(3), 647-665.

[22] Demsar, J. (2006). Statistical comparisons of classifiers over multiple data sets. *Journal of Machine Learning Research*, 7, 1-30.

[23] Friedman, M. (1937). The use of ranks to avoid the assumption of normality implicit in the analysis of variance. *Journal of the American Statistical Association*, 32(200), 675-701.

[24] Nemenyi, P. B. (1963). *Distribution-free Multiple Comparisons*. PhD Thesis, Princeton University.

[25] Wilcoxon, F. (1945). Individual comparisons by ranking methods. *Biometrics Bulletin*, 1(6), 80-83.
